



جامعة الجلالة  
GALALA UNIVERSITY

# Artificial Intelligence Science Program

Chapter 4: Learning from Examples

# Forms of Learning

- The technology of machine learning has become a standard part of software engineering.
- Any time you are building a software system, even if you don't think of it as an AI agent, components of the system can potentially be improved **with machine learning**.
  - For example, software to analyze images of galaxies with a machine-learned model.



# Forms of Learning

- An agent is **learning** if it improves its performance after making observations about the world.
- When the **agent is a computer**,
  - we call it **machine learning**: a computer observes some data, builds a model based on the data, and uses it.
- Any component of an agent program can be improved by machine learning.



# Types of learning

- In **supervised learning** the agent observes input-output pairs and learns a function that maps from input to output.
  - For example, the inputs could be camera images, each one is either “**bus**” or “**pedestrian**,” etc. An output like this is called a **label**.
- In **unsupervised learning** the agent learns patterns in the input without any explicit feedback.
  - The most common unsupervised learning task is **clustering**.
    - For example, when shown millions of images taken from the Internet, a computer vision system can identify a large cluster of similar images “cats.”



# Types of learning

- In **reinforcement learning** the agent learns from a series of reinforcements: **rewards** and **punishments**.
  - For example, at the end of a chess game the agent is told that it has won (a reward) or lost (a punishment).





# Supervised Learning

- Given a training set of example input–output pairs

$$(x_1, y_1), (x_2, y_2), \dots (x_N, y_N),$$

- where each pair was generated by an unknown function  $y = f(x)$   
discover a function that approximates the true function  $f$
- We can say  $y$  is the **ground truth**
- We can evaluate that with a second sample of  $(x_i, y_i)$  pairs called a **test set**.



# Prediction Problems: Classification vs. Numeric Prediction

- **Classification**
  - predicts categorical class **labels** (discrete or nominal)
  - classifies data (constructs a model) based on the **training set** and the values (**class labels**) in a classifying attribute and uses it in classifying new data.
- **Numeric Prediction**
  - models continuous-valued functions, i.e., predicts unknown or missing values
- **Typical applications**
  - Credit/loan approval:
  - Medical diagnosis: if a tumor is cancerous or benign.
  - Web page categorization: which category it is.



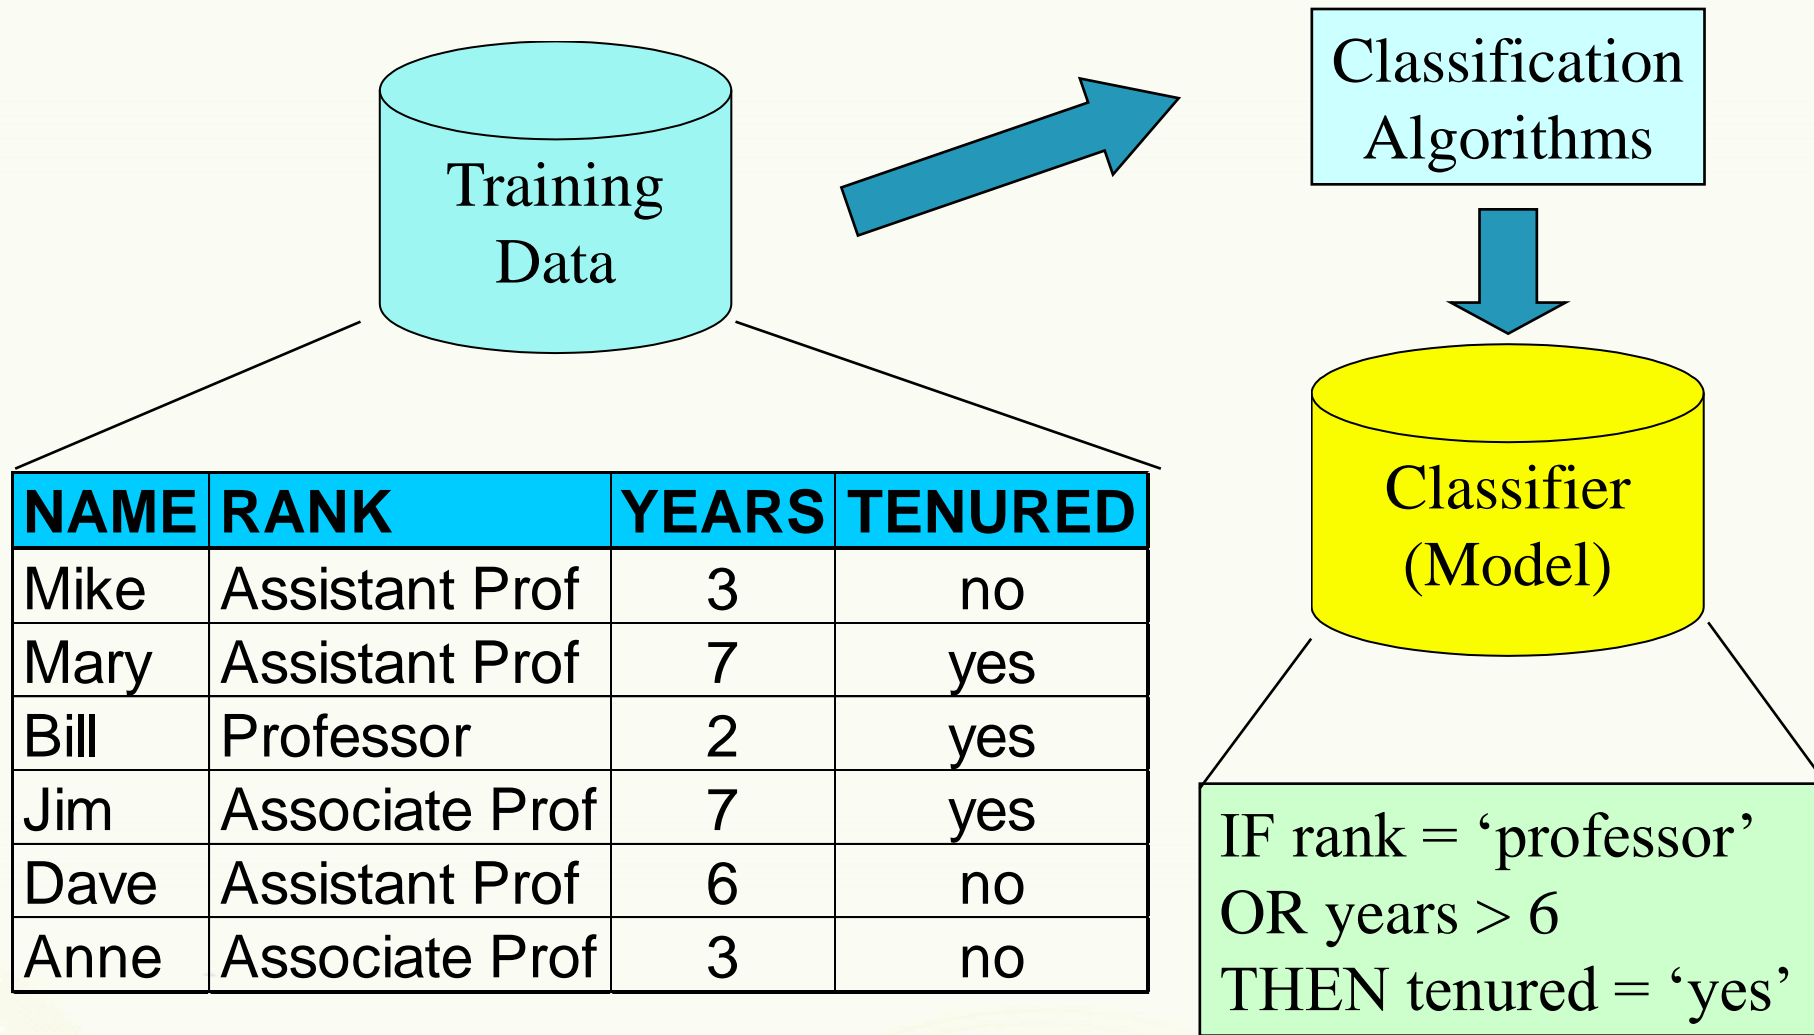
# Classification—A Two-Step Process

- **Model construction:** describing a set of predetermined classes
  - Each sample is assumed to belong to a predefined class, as determined by the **class label attribute**
  - The set of tuples used for model construction is **training set**
  - The model is represented as classification rules, decision trees, or mathematical formula
- **Model usage:** for classifying future or unknown objects
  - **Estimate accuracy** of the model
    - The known label of test sample is compared with the classified result from the model
    - **Accuracy** rate is the percentage of test set samples that are correctly classified by the model
    - **Test set** is independent of training set (otherwise overfitting)
  - If the accuracy is acceptable, use the model to **classify new data**
- Note: If *the test set* is used to select models, it is called **validation (test) set**

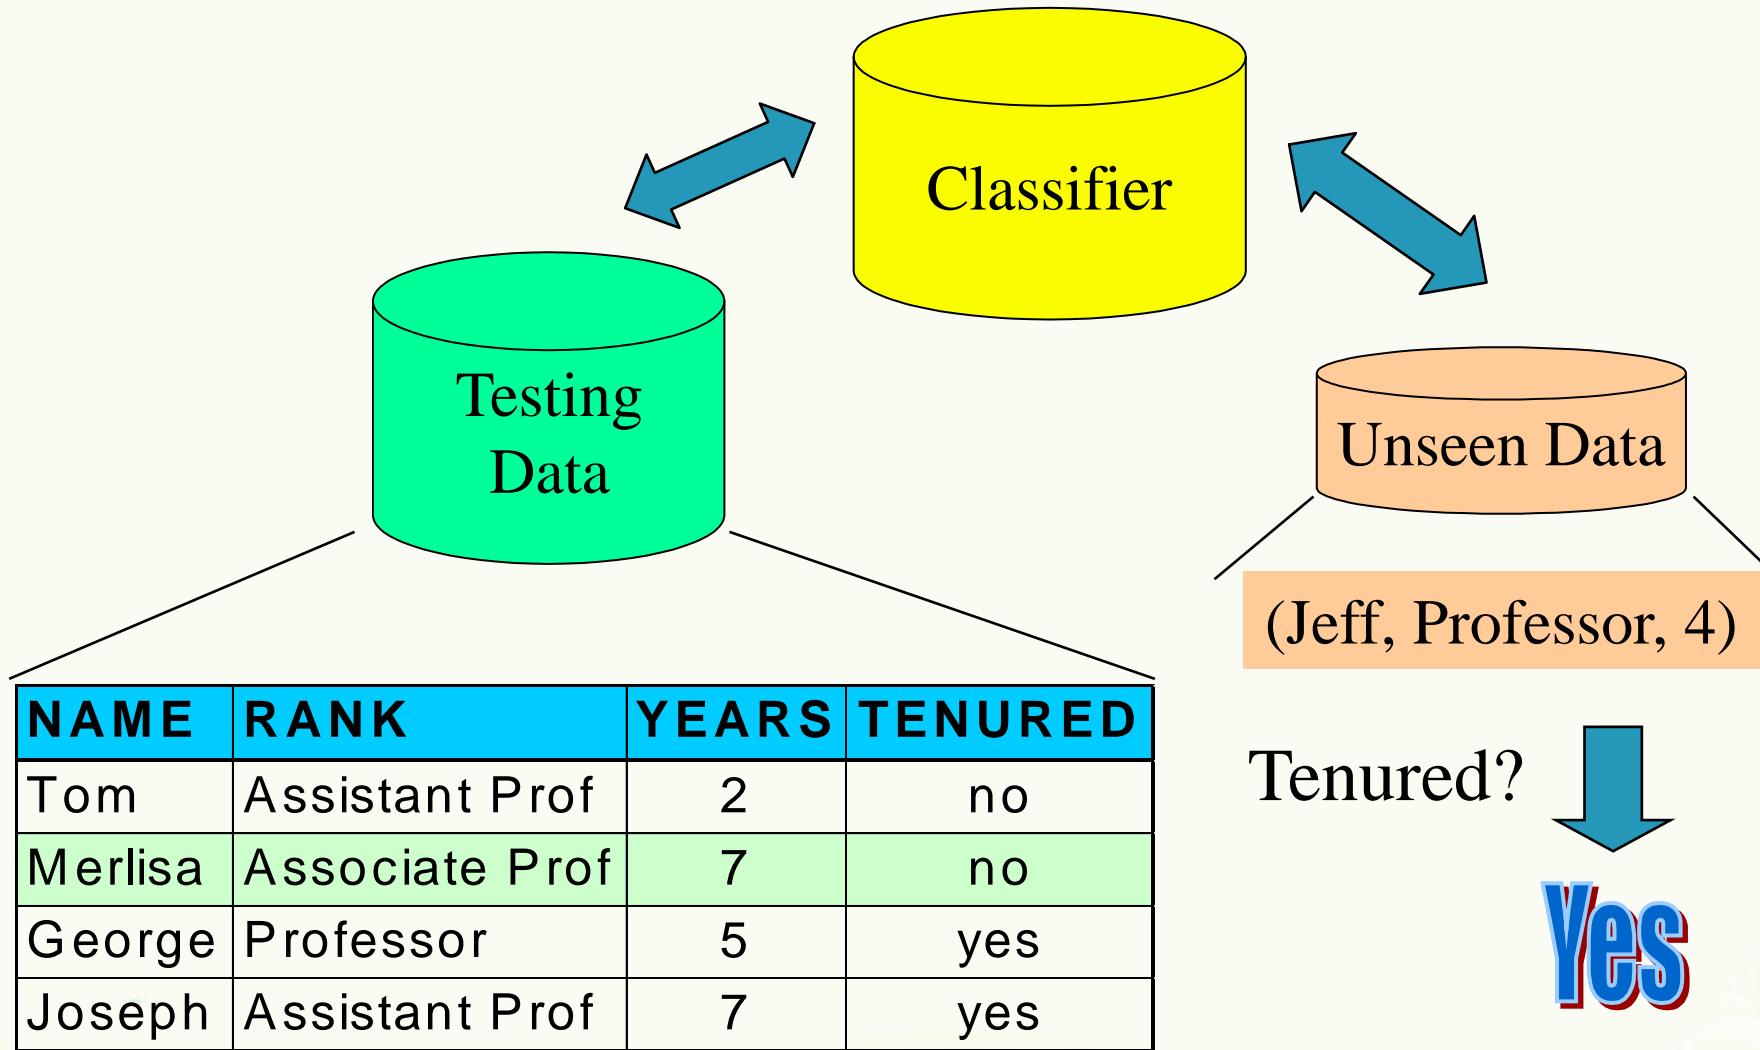




# Process (1): Model Construction



# Process (2): Using the Model in Prediction



# Learning Decision Trees

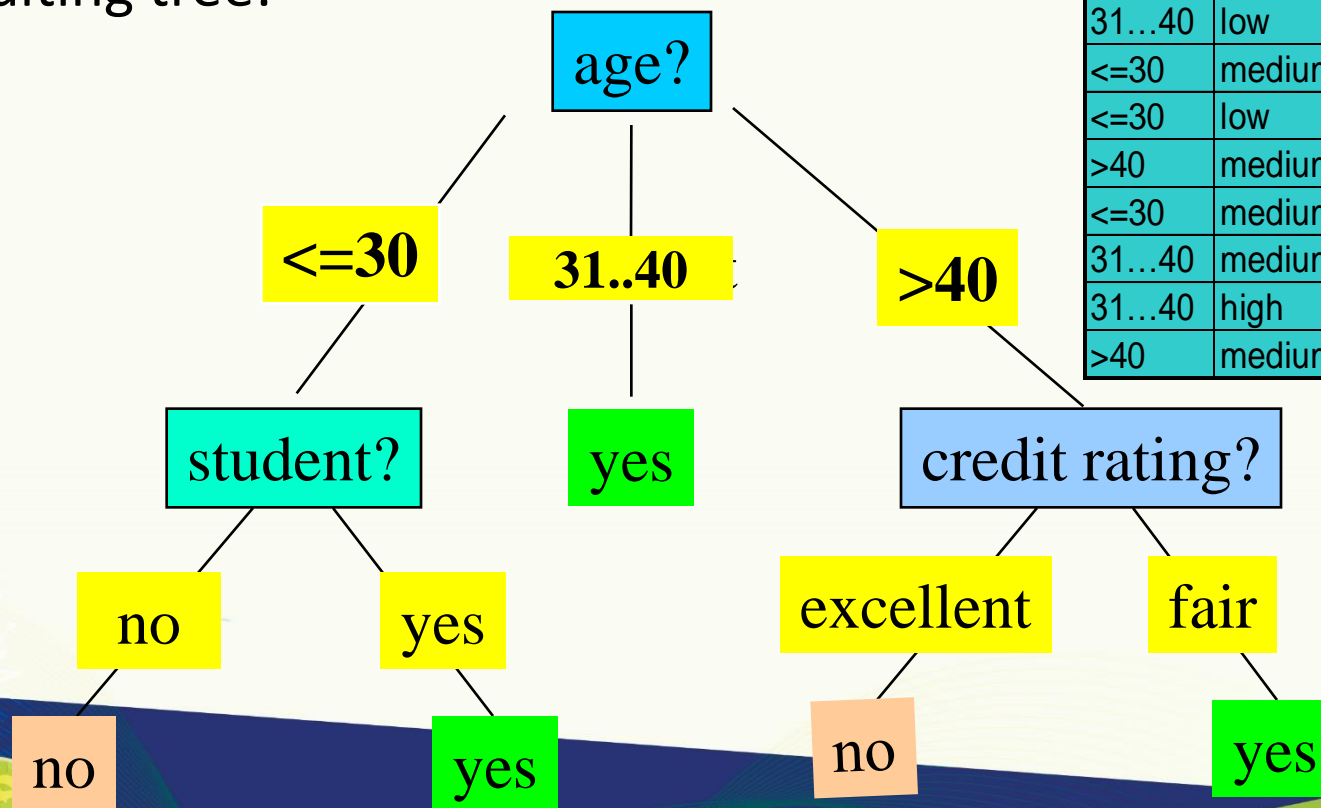
- A **decision tree** is a representation of a function that **maps a vector of attribute values** to a **single output value**—a “decision.”
- A decision tree reaches its decision by performing a sequence of tests, **starting at the root** and following the appropriate branch until a leaf is reached.



# Decision Tree Induction: An Example

- ❑ Training data set: Buys\_computer
- ❑ The data set follows an example of Quinlan's ID3 (Playing Tennis)
- ❑ Resulting tree:

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no



# Algorithm for Decision Tree Induction

- Basic algorithm (**a greedy algorithm**)
  - Tree is constructed in a **top-down recursive divide-and-conquer manner**.
  - At start, all the training examples are at the root
  - Attributes are categorical (if continuous-valued, they are discretized in advance)
  - Examples are partitioned recursively based on selected attributes
  - Test attributes are selected on the basis of a heuristic or statistical measure (e.g., **information gain**)
- Conditions for stopping partitioning
  - All samples for a given node belong to the same class
  - There are no remaining attributes for further partitioning – **majority voting** is employed for classifying the leaf
  - There are no samples left






# Brief Review of Entropy

- Entropy (Information Theory)

- A measure of uncertainty associated with a random variable
- Calculation: For a discrete random variable  $Y$  taking  $m$  distinct values  $\{y_1, \dots, y_m\}$ ,
  - $H(Y) = -\sum_{i=1}^m p_i \log(p_i)$ , where  $p_i = P(Y = y_i)$
- Interpretation:
  - Higher entropy => higher uncertainty
  - Lower entropy => lower uncertainty

- Conditional Entropy

- $H(Y|X) = \sum_x p(x) H(Y|X = x)$



$m = 2$



# Attribute Selection Measure: Information Gain (ID3/C4.5)

- Select the attribute with the highest information gain
- Let  $p_i$  be the probability that an arbitrary sample in  $D$  belongs to class  $C_i$ , estimated by  $|C_{i,D}|/|D|$ .

- Expected information (entropy) needed to classify a tuple in  $D$ :

$$Info(D) = -\sum_{i=1}^m p_i \log_2(p_i)$$

- Information needed (after using  $A$  to split  $D$  into  $v$  partitions) to classify  $D$ :

$$Info_A(D) = \sum_{j=1}^v \frac{|D_j|}{|D|} \times Info(D_j)$$

- Information gained by branching on attribute  $A$

$$Gain(A) = Info(D) - Info_A(D)$$



# Attribute Selection: Information Gain

■ Class P: buys\_computer = “yes”

■ Class N: buys\_computer = “no”

$$Info(D) = I(9,5) = -\frac{9}{14} \log_2\left(\frac{9}{14}\right) - \frac{5}{14} \log_2\left(\frac{5}{14}\right) = 0.940$$

age	p <sub>i</sub>	n <sub>i</sub>	I(p <sub>i</sub> , n <sub>i</sub> )
<=30	2	3	0.971
31...40	4	0	0
>40	3	2	0.971

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

$$Info_{age}(D) = \frac{5}{14} I(2,3) + \frac{4}{14} I(4,0) + \frac{5}{14} I(3,2) = 0.694$$

$\frac{5}{14} I(2,3)$  means “age <=30” has 5 out of 14 samples, with 2 yes’es and 3 no’s.

Hence

$$Gain(age) = Info(D) - Info_{age}(D) = 0.246$$

Similarly,

$$Gain(income) = 0.029$$

$$Gain(student) = 0.151$$

$$Gain(credit\_rating) = 0.048$$



## #Importing required libraries

```
import pandas as pd
import numpy as np
from sklearn.datasets import load_iris
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
data = load_iris() print('Classes to predict: ', data.target_names)
#Extracting data attributes
X = data.data
### Extracting target/ class labels
y = data.target
```

```
print('Number of examples in the data:', X.shape[0])
#First four rows in the variable 'X'
X[:4]
```

## #Output

```
Out: array([[5.1, 3.5, 1.4, 0.2],
           [4.9, 3. , 1.4, 0.2],
           [4.7, 3.2, 1.3, 0.2],
           [4.6, 3.1, 1.5, 0.2]])
```





```
#Using the train_test_split to create train and test sets.  
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state = 47, test_size = 0.25)  
#Importing the Decision tree classifier from the sklearn library.  
from sklearn.tree import DecisionTreeClassifier  
clf = DecisionTreeClassifier(criterion = 'entropy')
```

```
#Training the decision tree classifier.  
clf.fit(X_train, y_train)
```

```
#Output:
```

```
Out:DecisionTreeClassifier(class_weight=None, criterion='entropy', max_depth=None,  
    max_features=None, max_leaf_nodes=None,  
    min_impurity_decrease=0.0, min_impurity_split=None,  
    min_samples_leaf=1, min_samples_split=2,  
    min_weight_fraction_leaf=0.0, presort=False, random_state=None,  
    splitter='best')
```

```
#Predicting labels on the test set.  
y_pred = clf.predict(X_test)
```





```
#Importing the accuracy metric from sklearn.metrics library
```

```
from sklearn.metrics import accuracy_score
```

```
print('Accuracy Score on train data: ', accuracy_score(y_true=y_train, y_pred=clf.predict(X_train)))
```

```
print('Accuracy Score on test data: ', accuracy_score(y_true=y_test, y_pred=y_pred))
```

```
#Output:
```

```
Out: Accuracy Score on train data: 1.0
```

```
Accuracy Score on test data: 0.9473684210526315
```

